



CS 33: Introduction to Computer Organization

TA: Aalisha Dalal

LA: Jonathan Myong

Office Hours: Friday, 9:30-11:30AM

Outline



- **Parallel Computing/ High Performance Computing**
- **Using OpenMP**
- **Worksheet Problems**
- **Midterm Papers**

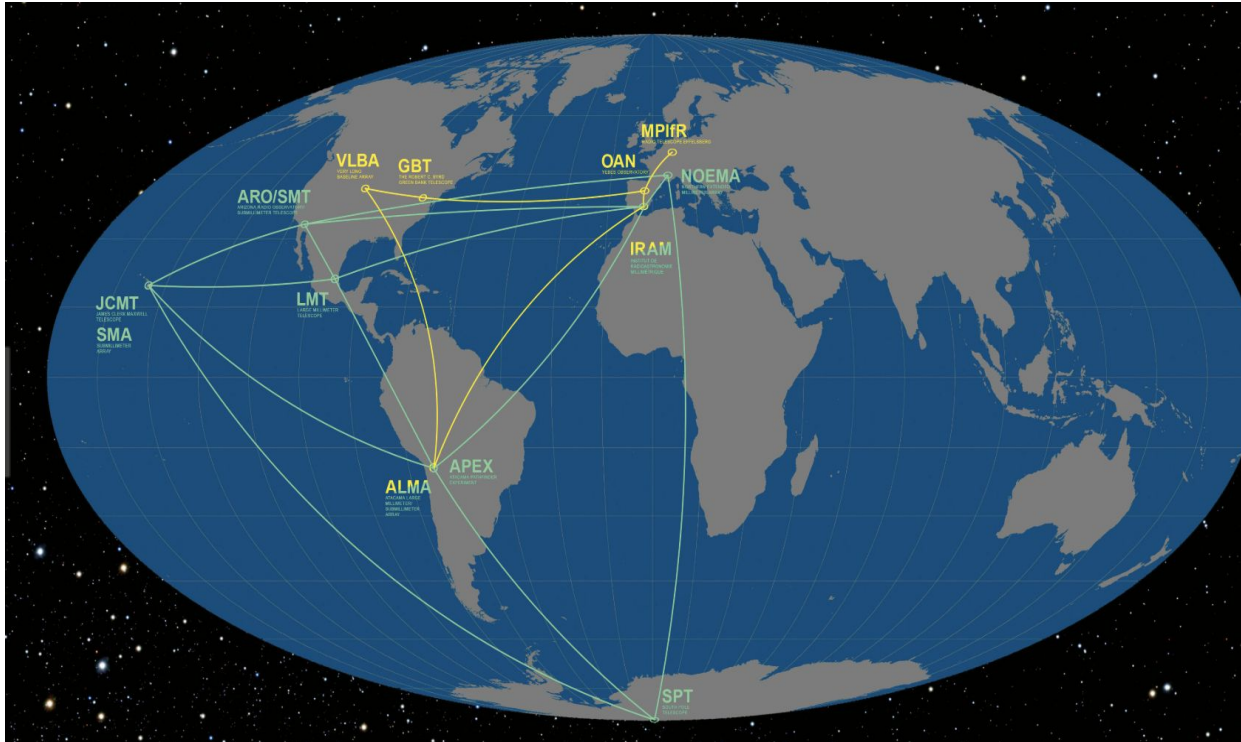
Why Parallel Computing?

“Faster computation by dividing the tasks among multiple processors instead of one.”

I wish there were more counters!



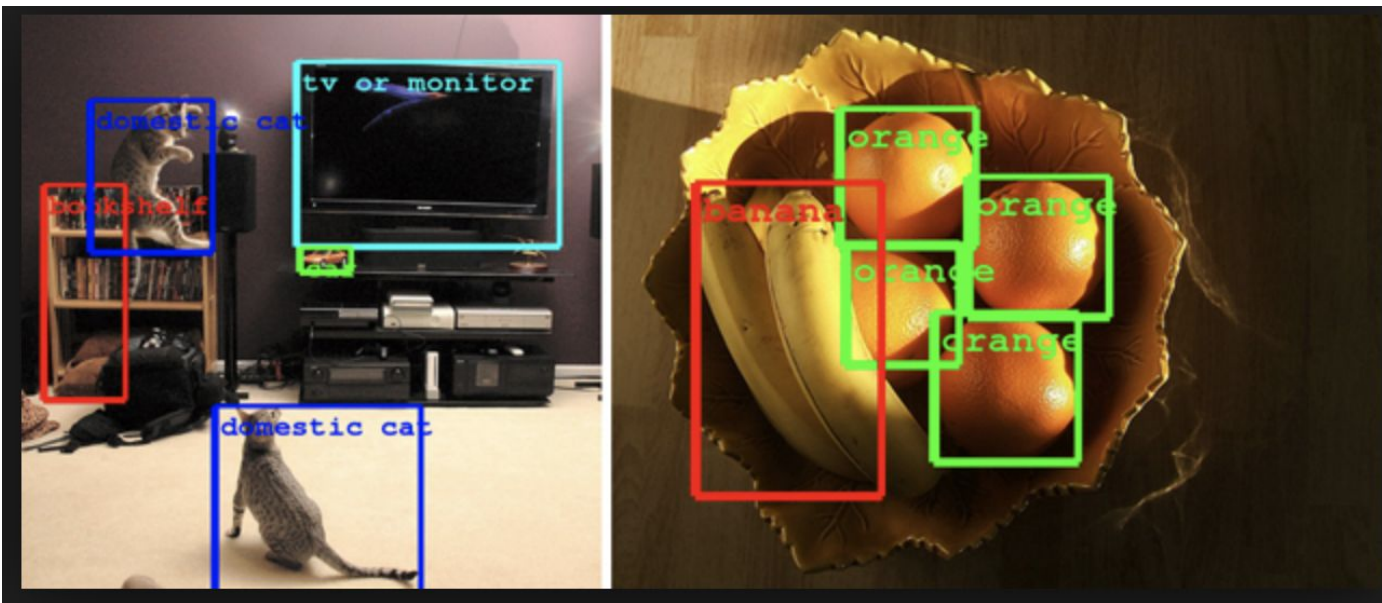
Motivation - To know our universe better!



Capturing the image
of the black hole!

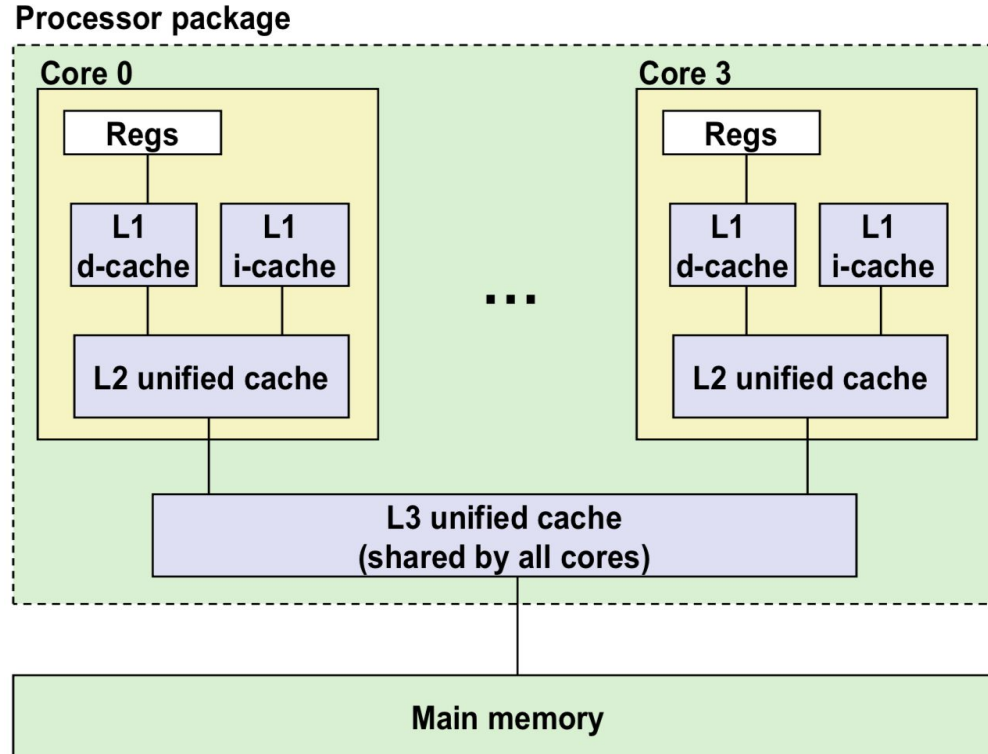


Motivation - In the Age of ML/AI



Task of Object
Detection using
CNN models and
many more

Parallel Architecture



Parallelism




- *Domain Decomposition*: dividing data among processors
- *Task Decomposition*: dividing the tasks/ operations among different processors
- *Pipelining*

Domain Decomposition

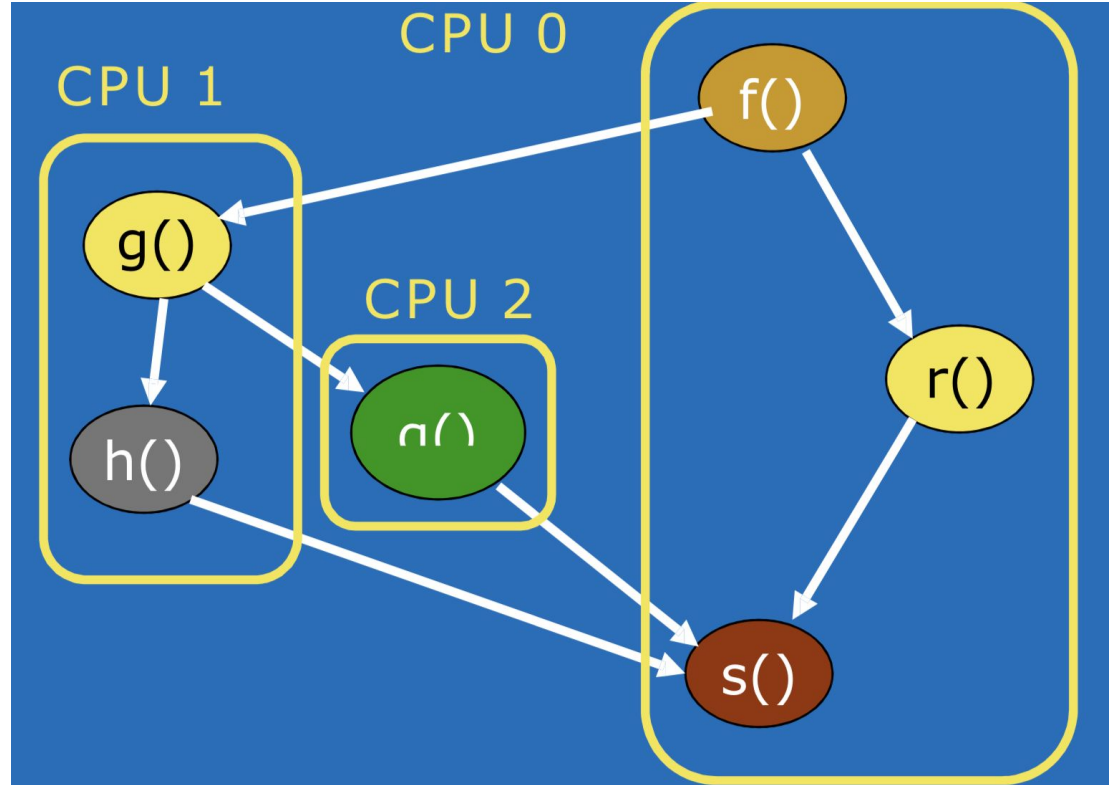


Finding the largest element in the array



Task Decomposition

Dividing the method calls among different CPU cores



Dependency structure

- *For parallelizing a code, it is crucial to understand the dependency graph structure.*

```
for (i = 0; i < 3; i++)  
    a[i] = b[i] / 2.0;
```


Loop iteration independent

Loop iteration dependent

```
for (i = 1; i < 4; i++)  
    a[i] = a[i-1] * b[i];
```

OpenMP

- API developed in C especially designed for programming using shared multi-processors
- Based on the ***Fork-Join*** Thread model




OpenMP Thread Creation

- Each thread executes the same code redundantly.

```
double A[1000];  
omp_set_num_threads(4)
```

```
double A[1000];  
omp_set_num_threads(4);  
#pragma omp parallel  
{  
    int ID = omp_get_thread_num();  
    pooh(ID, A);  
}  
printf("all done\n");
```

A single copy of A is shared between all threads.



Threads wait here for all threads to finish before proceeding (i.e. a *barrier*)

OpenMP Example

What does this
do?

```
#pragma omp parallel for
for (i=0; i < numPixels; i++)
{
    pGrayScaleBitmap[i] = (unsigned BYTE)
        (pRGBBitmap[i].red * 0.299 +
         pRGBBitmap[i].green * 0.587 +
         pRGBBitmap[i].blue * 0.114);
}
```

OpenMP Data Dependency Example



```
#pragma omp parallel for
for (i=2; i < 10; i++)      {
    factorial[i] = i * factorial[i-1];
}
```



*Data Dependency leads to inconsistency when
parallelizing code*


OpenMP Pragma Example

```
#pragma omp parallel
{
    #pragma omp for nowait
    for (i = 0; i < N; i++)
        a[i] = alpha(i);
    #pragma omp single nowait
    if (delta < 0.0) printf ("delta < 0.0\n");
    #pragma omp for
    for (i = 0; i < N; i++)
        b[i] = beta (i, delta);
}
```

What is the difference between “single” keyword, “serial code” and critical region?

OpenMP Example - Computing Pi

Numerical Integration



Mathematically, we know that:

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

We can approximate the integral as a sum of rectangles:

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$

Where each rectangle has width Δx and height $F(x_i)$ at the middle of interval i .

OpenMP Example - Computing Pi (contd.)

```
double area, pi, x;
int i, n;
...
area = 0.0;
for (i = 0; i < n; i++) {
    x = (i + 0.5) / n;
    area += 4.0 / (1.0 + x*x);
}
pi = area / n;
```

What would be the issue if we parallelize the “**for**” loop?

OpenMP Example - Computing Pi (contd.)

```
double area, pi, x;
int i, n;
...
area = 0.0;
for (i = 0; i < n; i++) {
    x = (i + 0.5) / n;
    area += 4.0 / (1.0 + x*x);
}
pi = area / n;
```

What would be the issue if we parallelize the “**for**” loop?

- RACE
CONDITION

Possible Solution

```
double area, pi, tmp, x;
int i, n;
...
area = 0.0;
#pragma omp parallel private(tmp)
{
    tmp = 0.0;
#pragma omp for private (x)
    for (i = 0; i < n; i++) {
        x = (i + 0.5)/n;
        tmp += 4.0/(1.0 + x*x);
    }
#pragma omp critical
    area += tmp;
}
pi = area / n;
```

What did we do and
why is this better?

Possible Solution - Reduction operator



```
double area, pi, x;
int i, n;
...
area = 0.0;
#pragma omp parallel for private(x) \
                        reduction(+:area)
for (i = 0; i < n; i++) {
    x = (i + 0.5)/n;
    area += 4.0/(1.0 + x*x);
}
pi = area / n;
```

This is a powerful
technique!

Locking Mechanism

What happens if
threads are at
this point at the
same time?

Thread A

```
lock (lock_a);
```

```
a += 5;
```

```
lock (lock_b);
```

```
b += 7;
```

```
a += b;
```

```
unlock (lock_b);
```

```
a += 11;
```

```
unlock (lock_a);
```

Thread B

```
lock (lock_b);
```

```
b += 5;
```

```
lock (lock_a);
```

```
a += 7;
```

```
a += b;
```

```
unlock (lock_a);
```

```
b += 11;
```

```
unlock (lock_b);
```

If there is an issue,
how can we solve
it?



Worksheet

<https://tinyurl.com/cs33-parallel-universe>



MidTerm Papers



Thank you!

Appendix



Threads vs. Processes

■ Threads and processes: similarities

- Each has its own logical control flow
- Each can run concurrently with others
- Each is context switched (scheduled) by the kernel

■ Threads and processes: differences

- Threads share code and data, processes (typically) do not
- Threads are much less expensive than processes
 - Process control (creating and reaping) is more expensive as thread control
 - Context switches for processes much more expensive than for threads